

# Resource Provisioning with using Deep Reinforcement Learning

Baochun Li

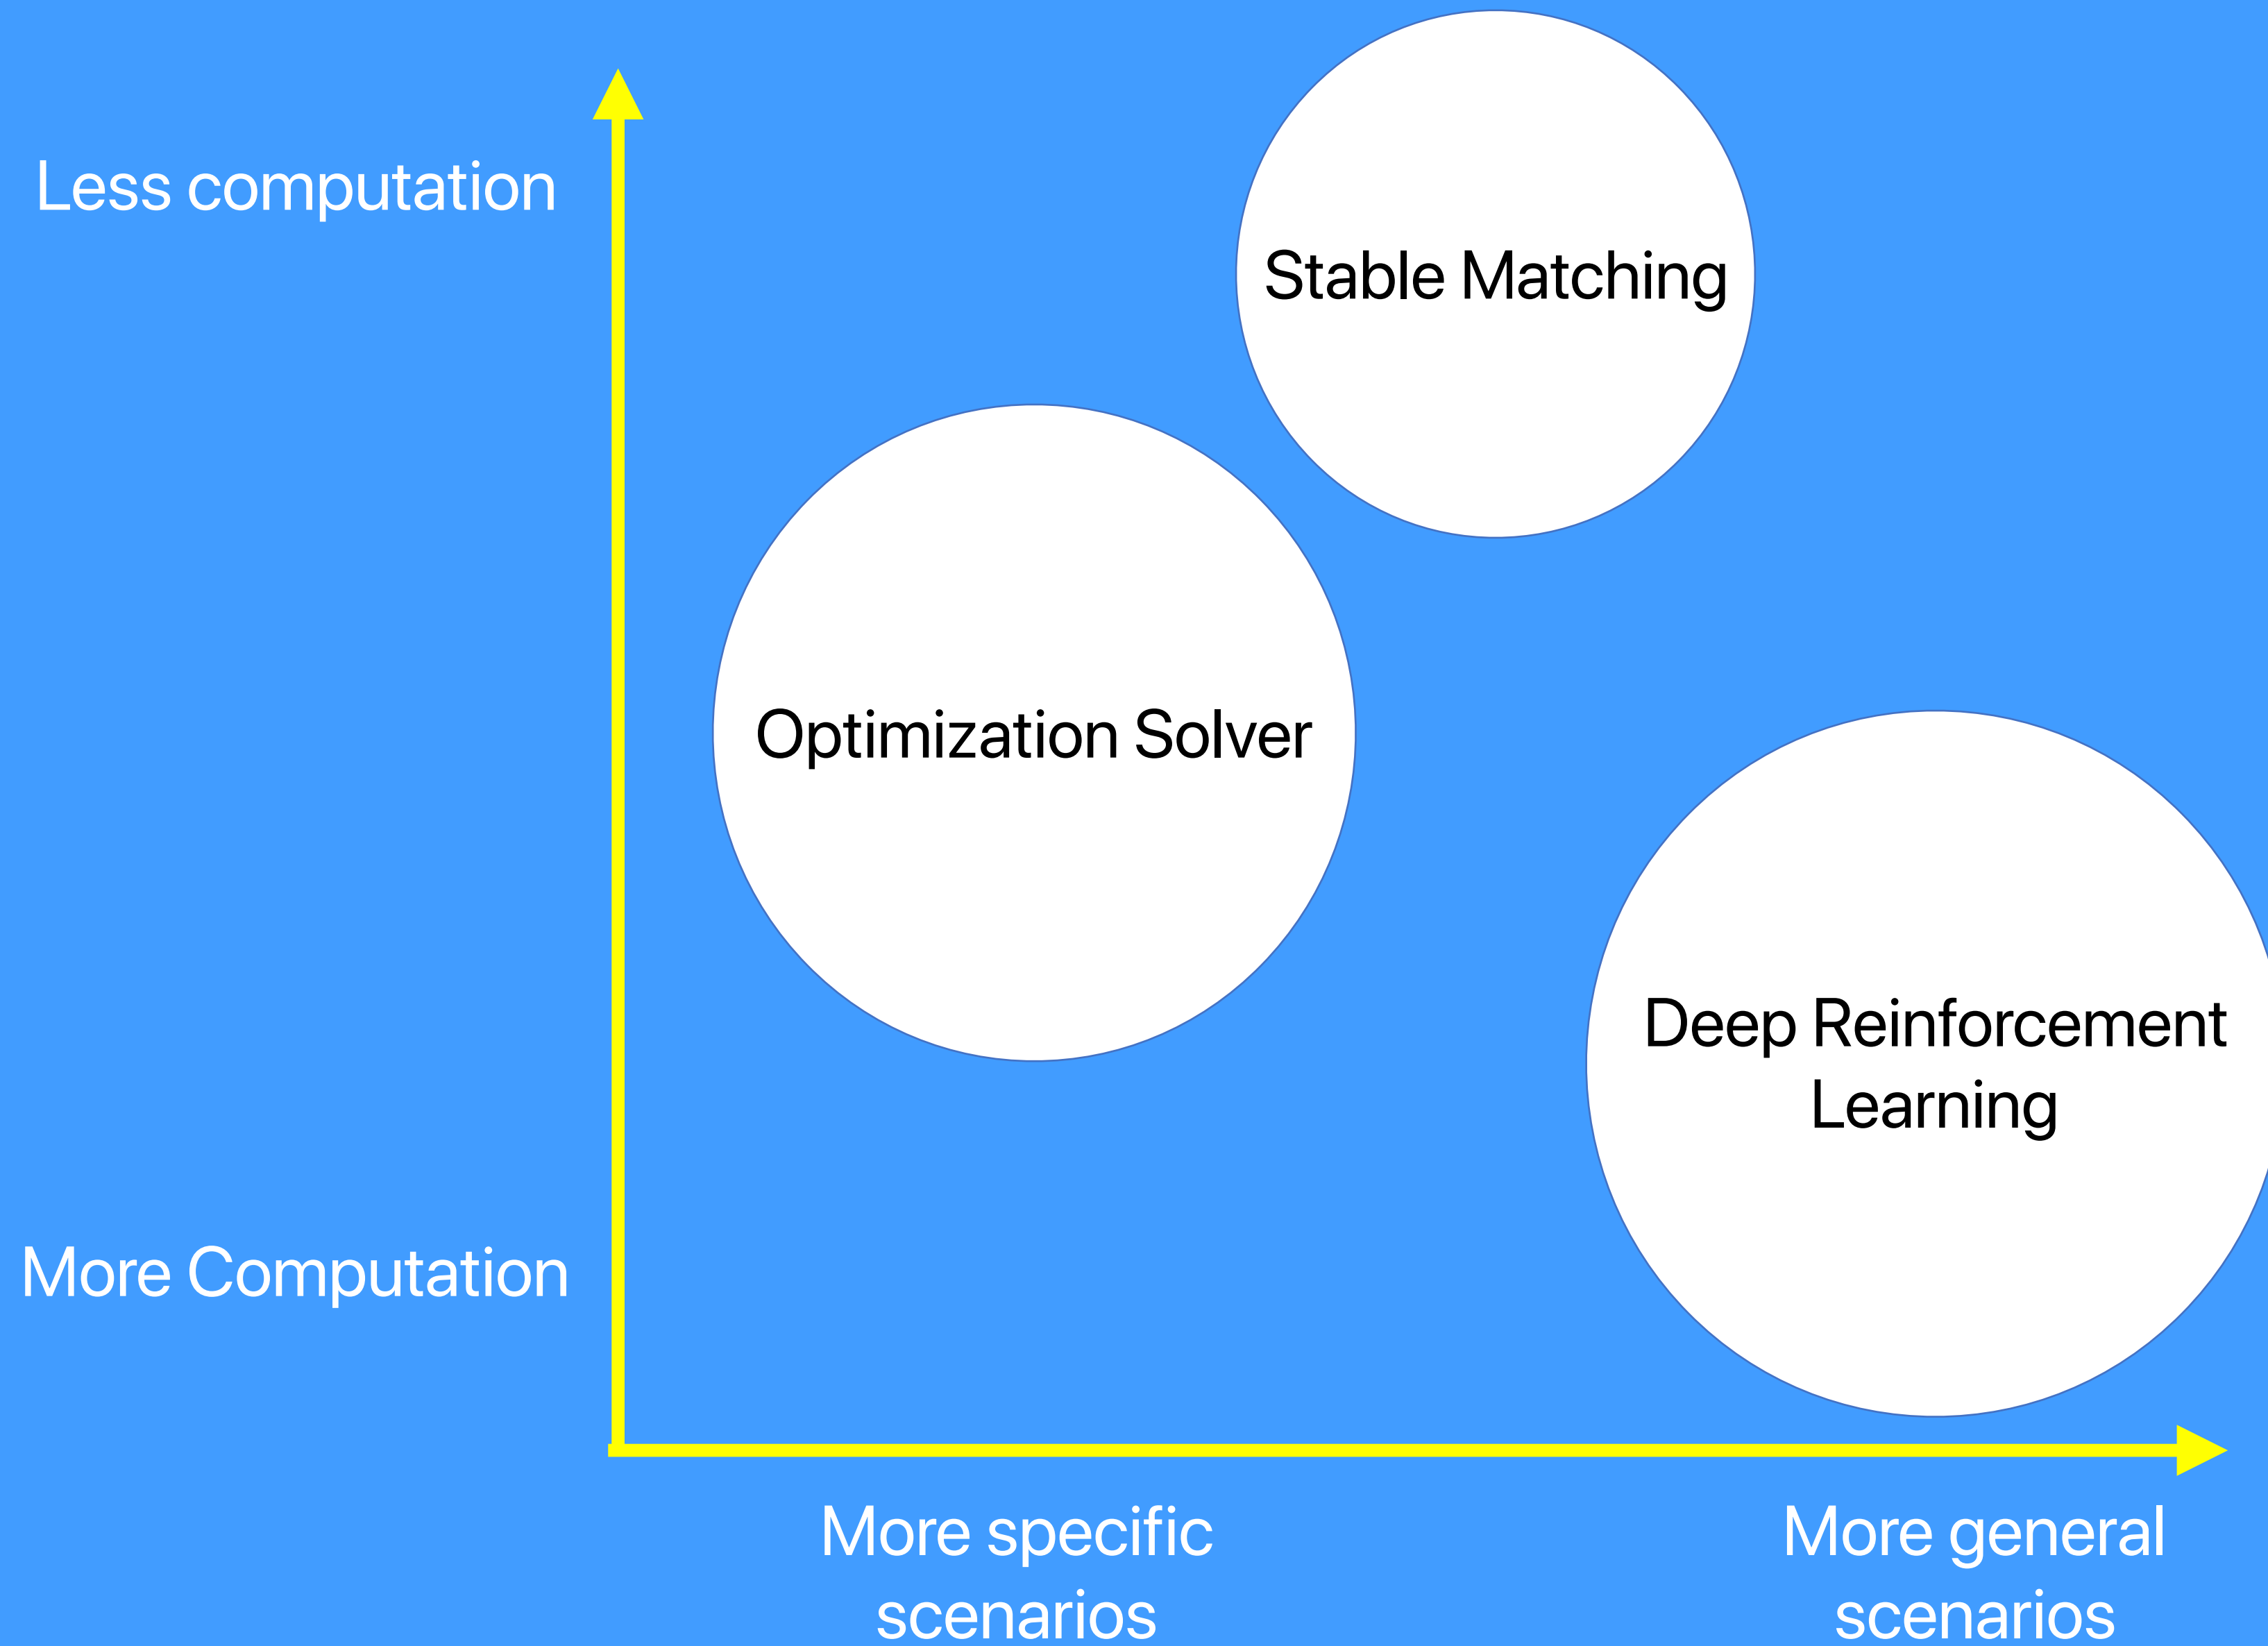
Department of Electrical and Computer Engineering

University of Toronto

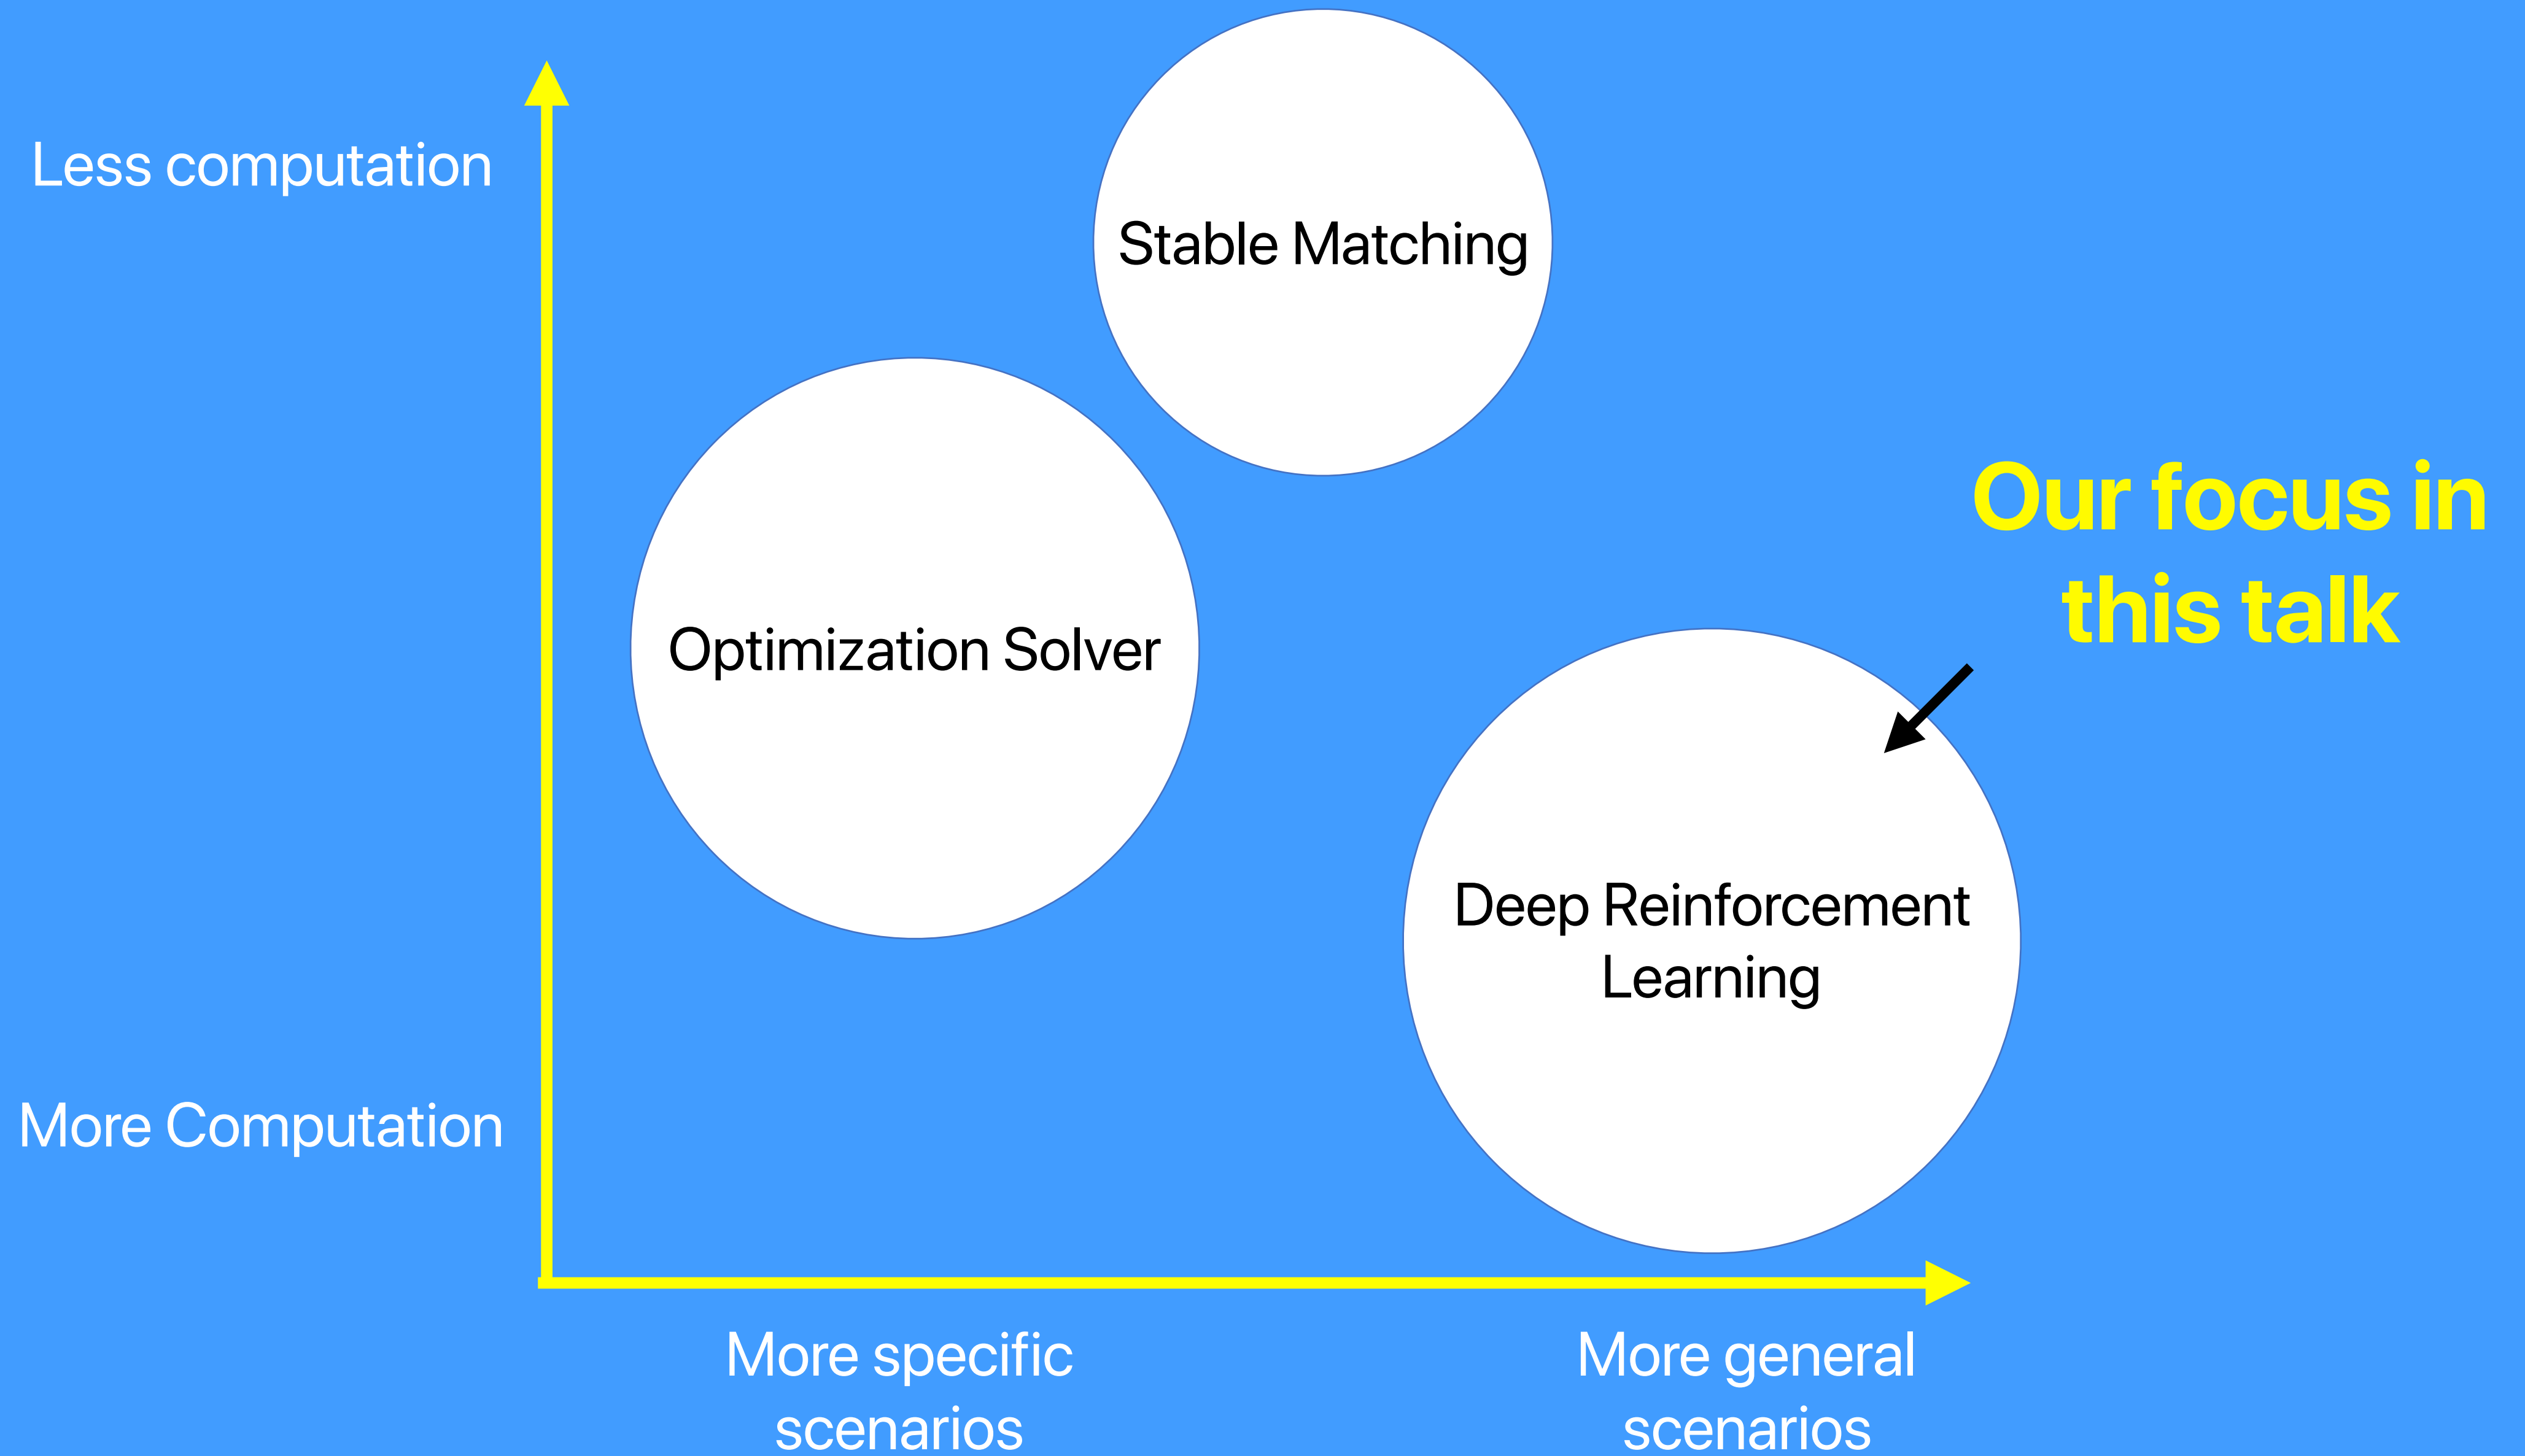
IWQoS 2019, Phoenix, Arizona

June 24, 2019

**Resource assignment problems in  
general: minimize the job completion time  
by **assigning resources****



Solving **resource assignment** problems



Solving **resource assignment** problems


**One running example today**

Modern machine learning training workloads need a large amount of computation resources, and can take **days** to complete

**Objective:** complete the machine learning training workload as soon as possible

**How?** By optimally assigning devices — GPUs and CPUs — to neural network operations

# Neural networks: computation graph





# Same neural network, but in Python code

```
# activation function (use sigmoid)
f = lambda x: 1.0/(1.0 + np.exp(-x))

# random input vector of three numbers (3x1)
x = np.random.randn(3, 1)

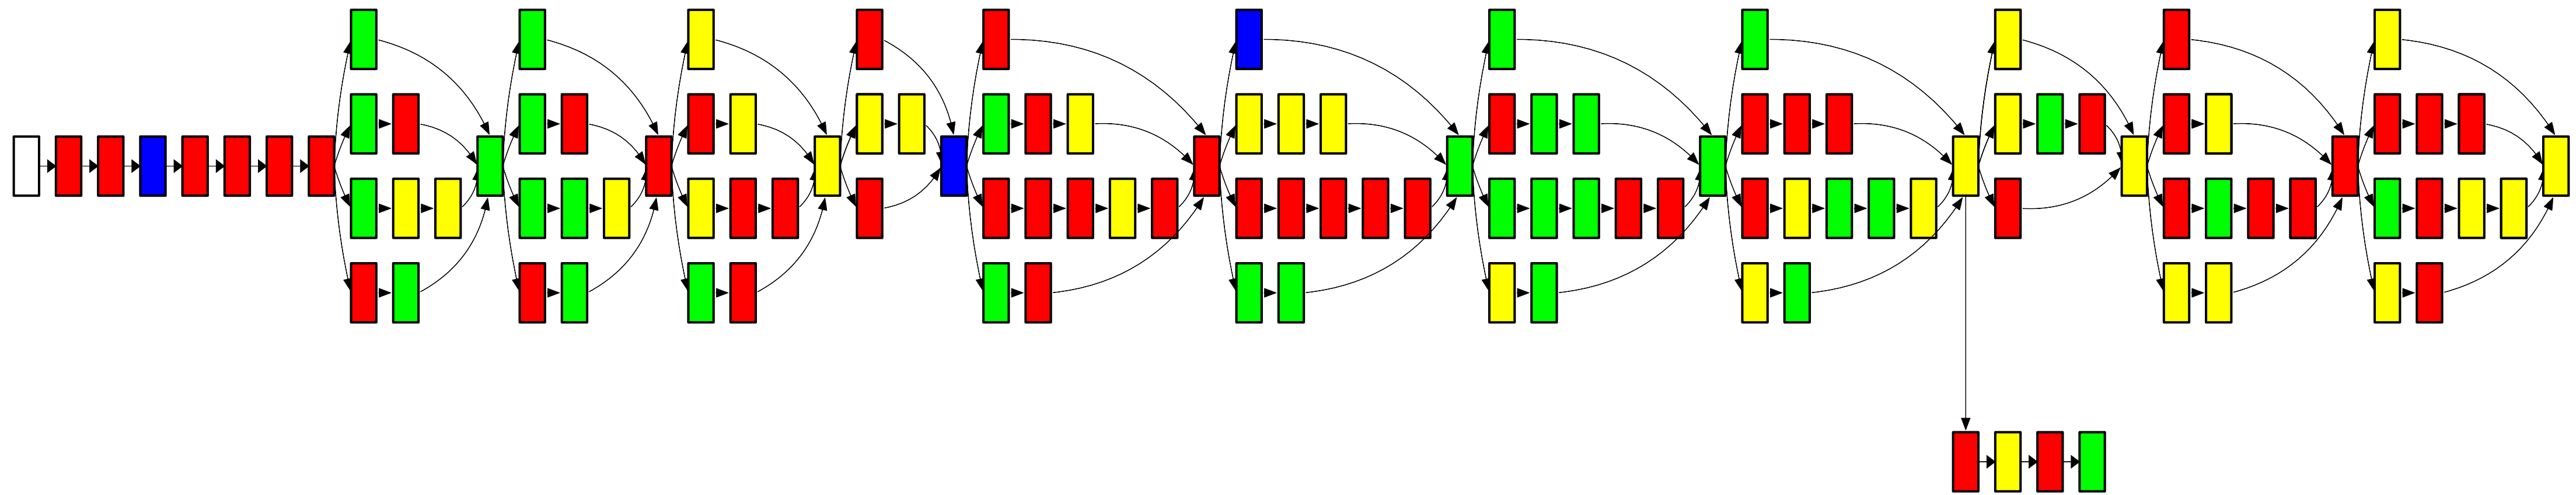
# W1, W2, W3, b1, b2, b3 are learnable parameters
# calculate first hidden layer activations (4x1)
h1 = f(np.dot(W1, x) + b1)

# calculate second hidden layer activations (4x1)
h2 = f(np.dot(W2, h1) + b2)

# output neuron (1x1)
out = np.dot(W3, h2) + b3
```

# Device placement with deep reinforcement learning

**Objective:** to find the **best** way to assign devices to operations to minimize training time



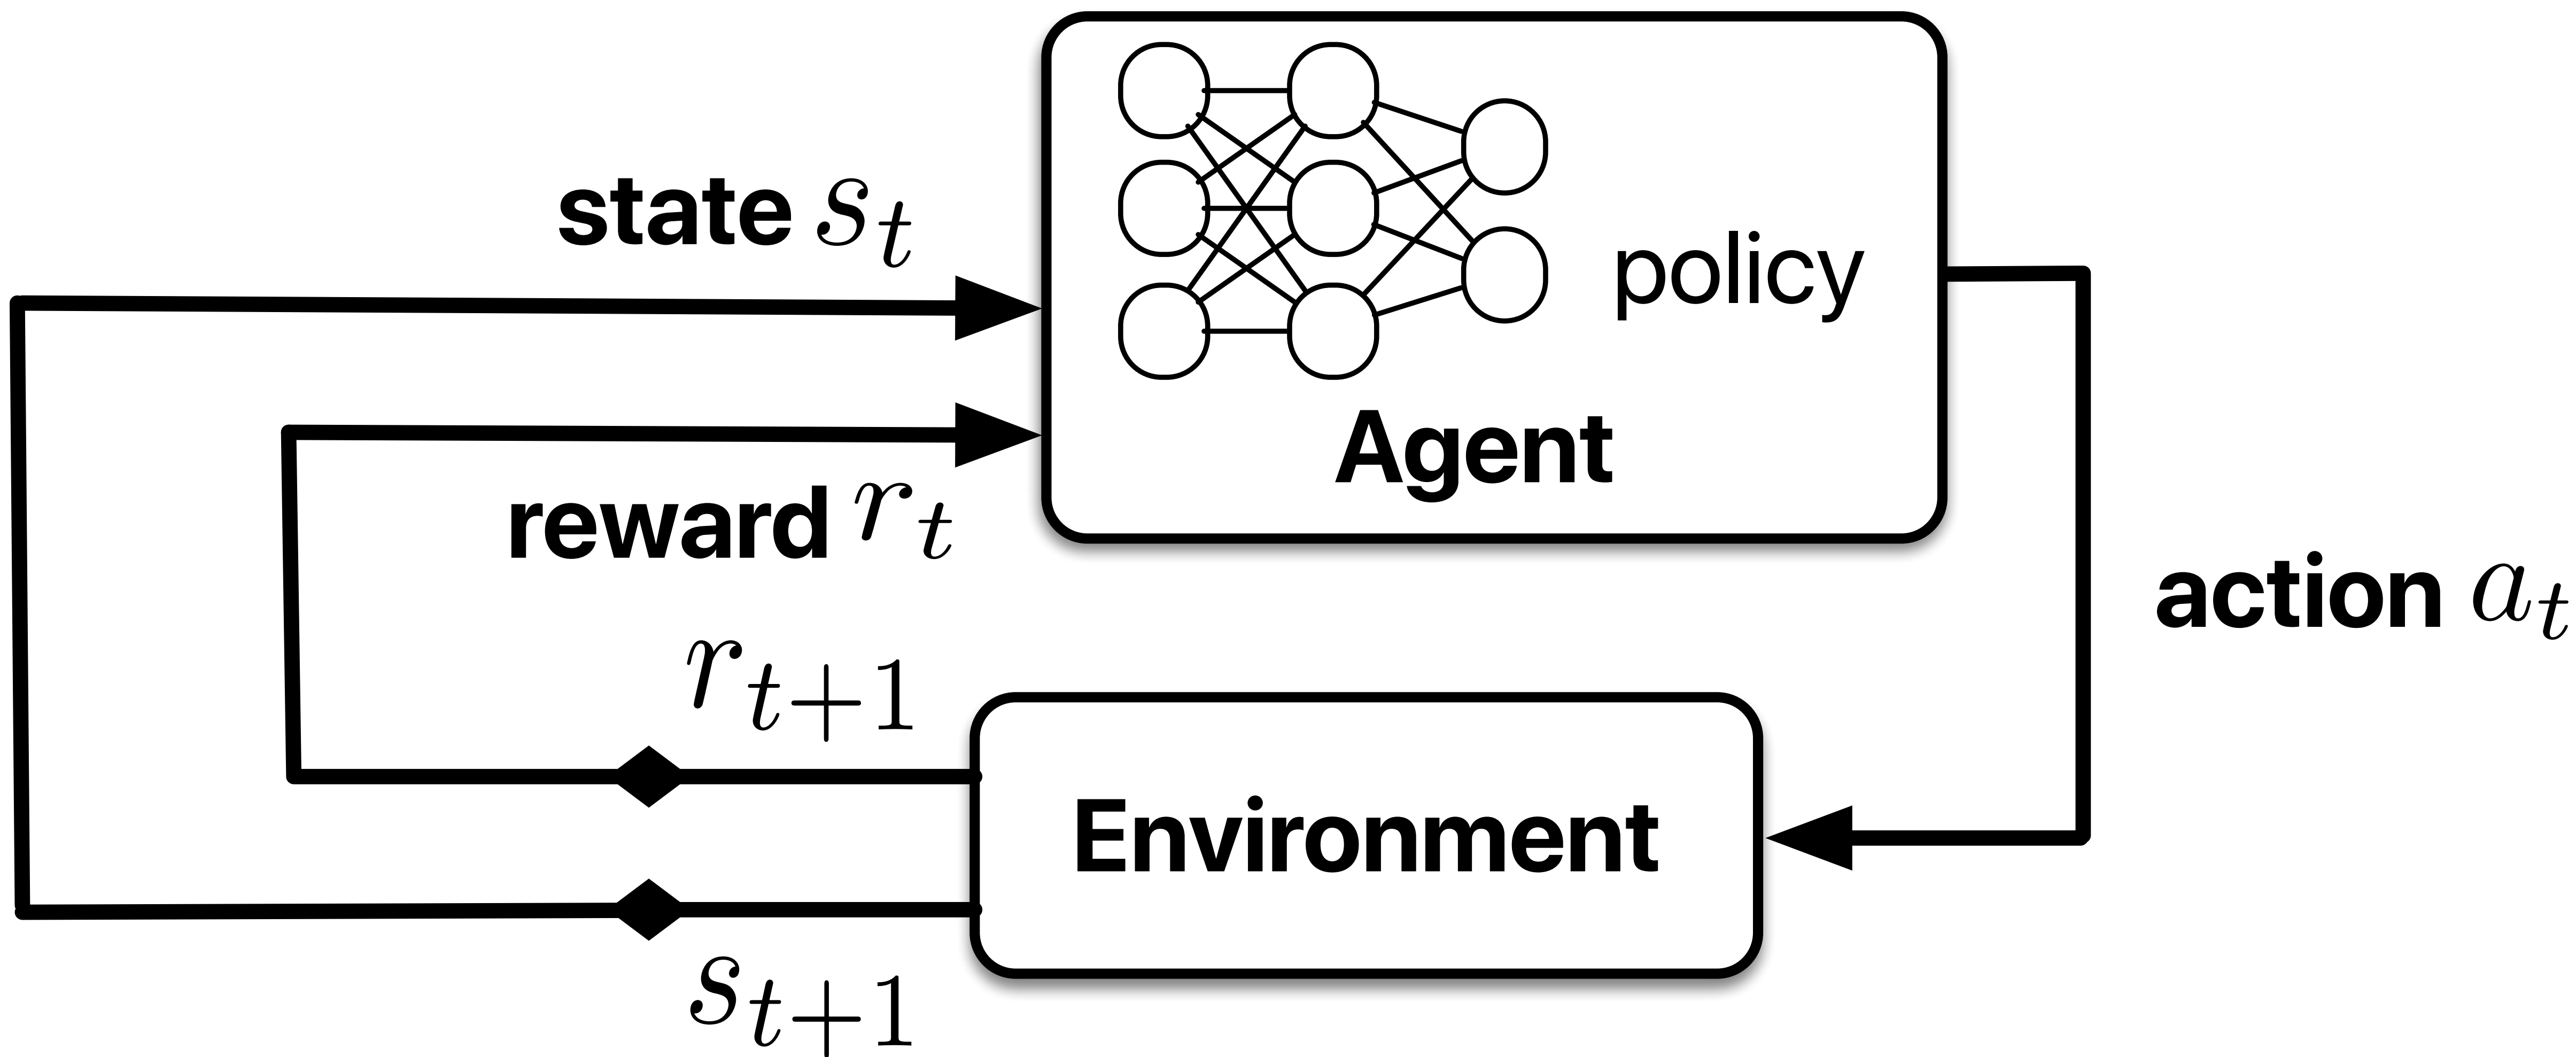
Mirhoseini et al. (Google Inc.), “**Device placement optimization with reinforcement learning**,” in Proc. **ICML 2017**.

**But what is deep reinforcement learning?**

# Recent success stories from DeepMind

- ▶ Deep Q-Network: Atari 2600 games (February 2015)
- ▶ AlphaGo (3:1 win over Lee Sedol, October 2016)
- ▶ AlphaGo Zero (100:0 win over AlphaGo, October 2017)
  - ▶ Learned from scratch using self-play with deep reinforcement learning and Monte-Carlo Tree Search

**Let's start from the beginning**




**Reinforcement learning is “semi-supervised,” with not much more guidance than a **reward** from the environment**



# Toy example: Cartpole

**States observed:** [position of cart, velocity of cart, angle of pole, pole velocity at tip]

**Possible actions:** [push cart to the left, push cart to the right]





# Cartpole: random agent

```
import gym
env = gym.make("CartPole-v0")
total_reward = 0.0
total_steps = 0
obs = env.reset()

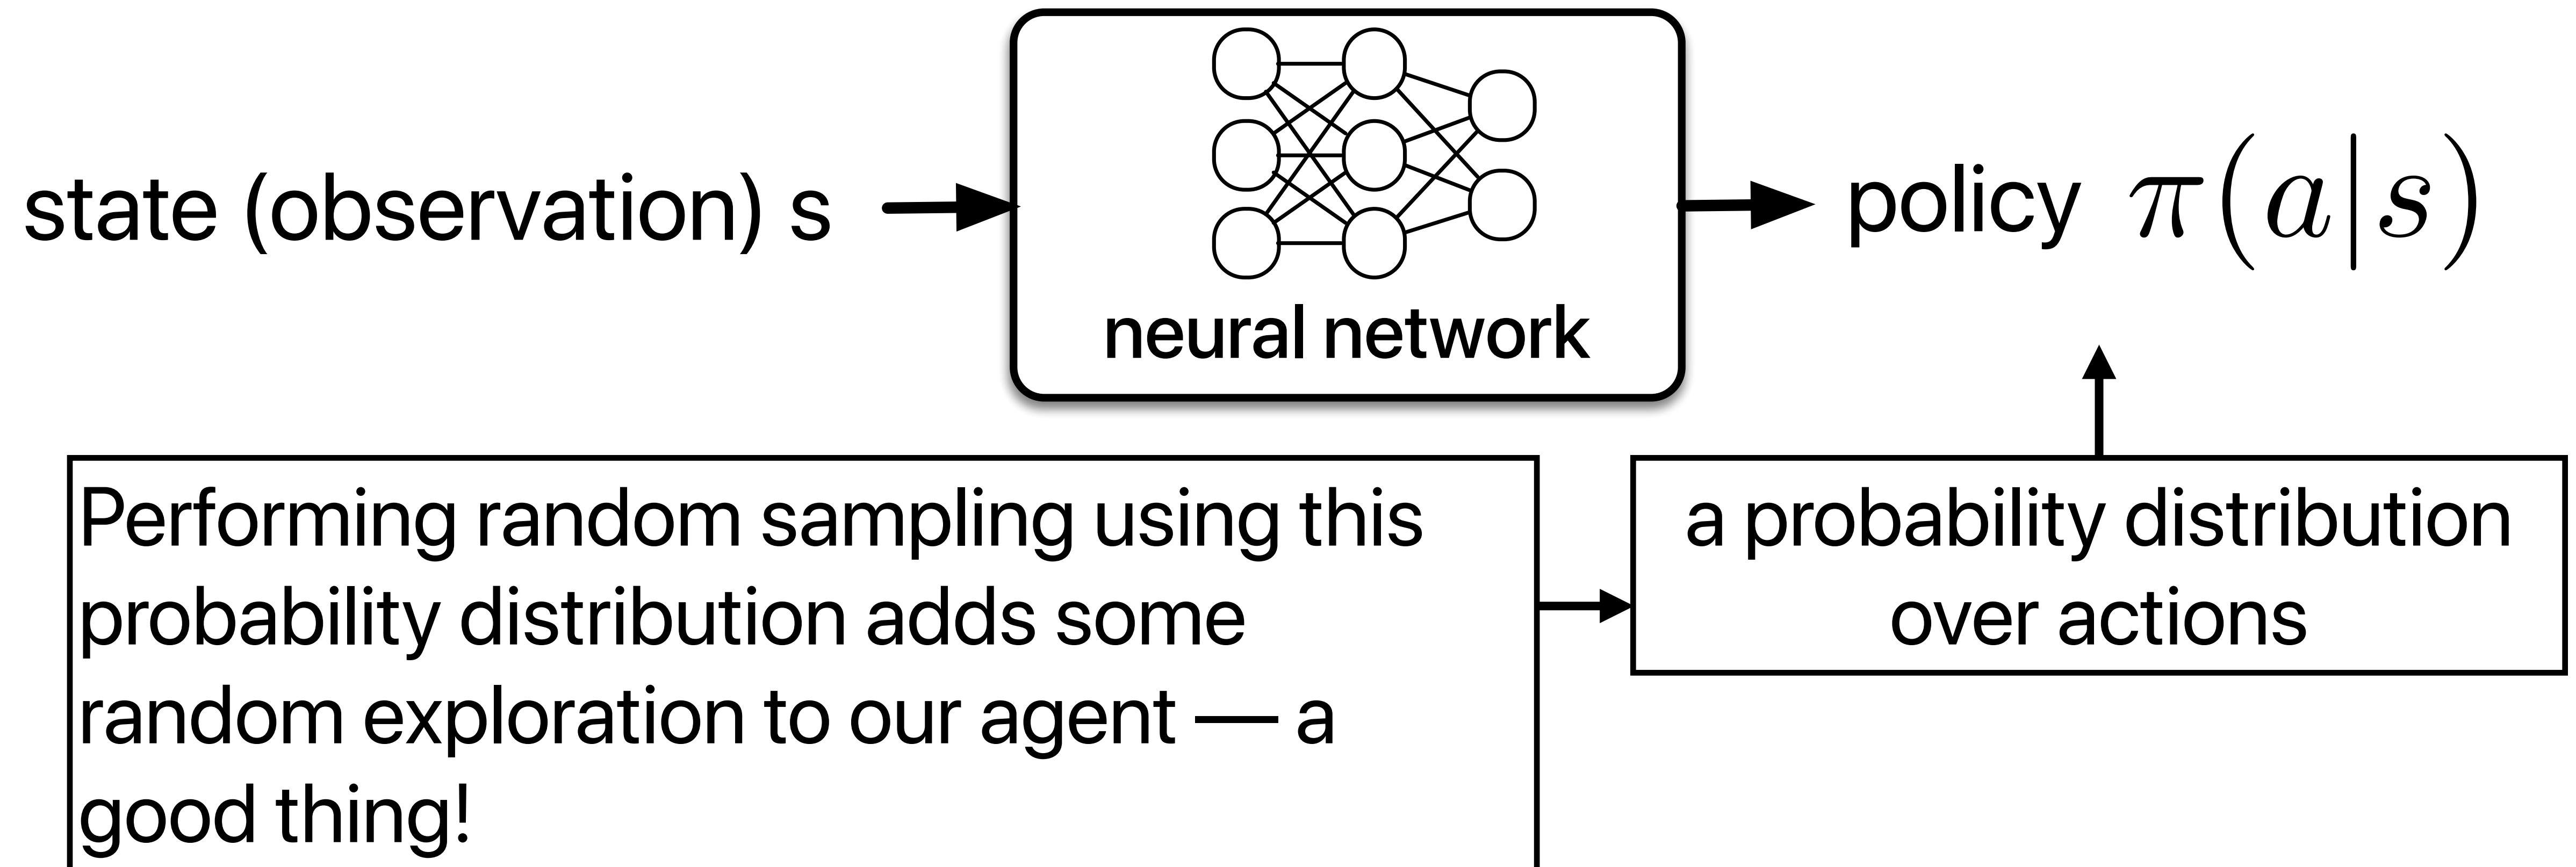
# start taking random actions
while True:
    action = env.action_space.sample()
    observation, reward, done, _ = env.step(action)
    total_reward += reward
    total_steps += 1
    if done:
        break
```

```
(pytorch) $ python cartpole_random.py
```

# **Our first algorithm: the cross-entropy method**

**Model free:** directly connects observations to actions, without building a model of the environment or the reward

# Our first algorithm: the cross-entropy method



# Cross-entropy: training

- ▶ Play N episodes using our current model and environment
- ▶ Calculate the total reward for every episode and decide on a reward boundary (say, 70%)
- ▶ Throw away all episodes with a reward below the boundary
- ▶ Train on the remaining **elite** episodes
- ▶ Repeat until we are satisfied with the result

# Cross-entropy: an iterative algorithm

$$\pi_{i+1}(a|s) = \arg \min_{\pi_{i+1}} - \mathbb{E}_{z \sim \pi_i(a|s)} [R(z) \geq \phi_i] \log \pi_{i+1}(a|s)$$



**Basic idea:** we sample episodes using our current policy (starting with some random initial policy), and minimize the **negative log-likelihood** of the most successful samples using our policy

- ▶ happens to be the same as minimizing the **cross-entropy** (and the **Kullback-Leibler (K-L) divergence** that quantifies the distance between two probability distributions)

# Defining our neural network using PyTorch

```
import torch.nn as nn

class Net(nn.Module):
    def __init__(self, obs_size, hidden_size, n_actions):
        super(Net, self).__init__()
        self.net = nn.Sequential(
            nn.Linear(obs_size, hidden_size),
            nn.ReLU(),
            nn.Linear(hidden_size, n_actions)
        )

    def forward(self, x):
        return self.net(x)
```

# Starting a loop to interact with the environment

```
while True:
    obs_v = torch.FloatTensor([obs])
    act_probs_v = sm(net(obs_v))
    act_probs = act_probs_v.data.numpy()[0]

    action = np.random.choice(len(act_probs),
                              p=act_probs)

    next_obs, reward, is_done, _ = env.step(action)
    episode_reward += reward
```



# Training our neural network

```
env = gym.make("CartPole-v0")
obs_size = env.observation_space.shape[0]
n_actions = env.action_space.n

net = Net(obs_size, HIDDEN_SIZE, n_actions)
objective = nn.CrossEntropyLoss()
optimizer = optim.Adam(params=net.parameters(), lr=0.01)

for iter_no, batch in enumerate(iterate_batches(env, net,
BATCH_SIZE)):
    obs_v, acts_v, reward_b, reward_m = filter_batch(batch, PERCENTILE)
    optimizer.zero_grad()
    action_scores_v = net(obs_v)
    loss_v = objective(action_scores_v, acts_v)
    loss_v.backward()
    optimizer.step()
```

combines **softmax** (exp) and **cross-entropy** (log) into one function for numerical stability

```
(pytorch) $ python cartpole_cross_entropy.py
```

# Our second algorithm: policy gradient

- ▶ A policy gradient method, called REINFORCE, was used by Google's ICML 2017 paper to solve the device placement problem
- ▶ The cross-entropy method uses the **elite** episodes with high rewards, and discards the bad episodes with low rewards
- ▶ But can we use a more **fine-grained separation** of episodes?
  - ▶ Perhaps an episode with a total reward of 100 should contribute more than another episode with a total reward of 50?

# REINFORCE: Monte-Carlo Policy Gradient

- ▶ Play N episodes using our current model and environment
- ▶ For every step  $t$  in every episode  $k$ , calculate a discounted total reward for subsequent steps:

$$Q_{k,t} \leftarrow \sum_{k=t+1}^T \gamma^{k-t-1} R_k$$

- ▶ An episode with rewards  $[1, 1, 1, 1]$  now becomes  $[3.9404, 2.9701, 1.99, 1.0]$  ( $\gamma = 0.99$ )

# REINFORCE: Monte-Carlo Policy Gradient

Update the neural network weights to minimize the loss function:

$$-\mathbb{E}_{k,t} \left[ \underbrace{Q_{k,t}}_{\text{discounted total reward}} \log \underbrace{\pi(a_{k,t} | s_{k,t}; \theta)}_{\text{policy function}} \right]$$

# Training our neural network

```
optimizer.zero_grad()
states_v = torch.FloatTensor(batch_states)
batch_actions_t = torch.LongTensor(batch_actions)
batch_qvals_v = torch.FloatTensor(batch_qvals)

logits_v = net(states_v)
log_prob_v = F.log_softmax(logits_v,
                             dim=1)
log_prob_actions_v = batch_qvals_v *
    log_prob_v[range(len(batch_actions)), batch_actions_t]
loss_v = -log_prob_actions_v.mean()
loss_v.backward()
optimizer.step()
```

← combines **softmax** (exp) and **log** into one function for numerical stability

```
(pytorch) $ python cartpole_reinforce.py
```

# We are not done yet: Proximal Policy Optimization

- ▶ Motivation: improve the **stability** of policy updates during training
  - ▶ We wish to train as fast as we can, making large steps in Stochastic Gradient Descent (SGD) updates
  - ▶ But since our policy is very nonlinear, a large update can ruin the policy we've just learned — making a bad update once may not be recoverable later
  - ▶ One can make tiny steps in SGD updates, but this will slow down convergence



**Lesson learned:** distance in parameter space  $\neq$  distance in policy space


# Proximal Policy Optimization

Update the neural network weights to minimize the loss function:

$$-\mathbb{E}_{k,t} \left[ Q_{k,t} \frac{\pi(a_{k,t} | s_{k,t}; \theta)}{\pi(a_{k,t} | s_{k,t}; \theta_{\text{old}})} \right] - \beta D_{\text{KL}}[\pi(\cdot | s_{k,t}, \theta) \parallel \pi(\cdot | s_{k,t}, \theta_{\text{old}})]$$

policy parameters  
before the update

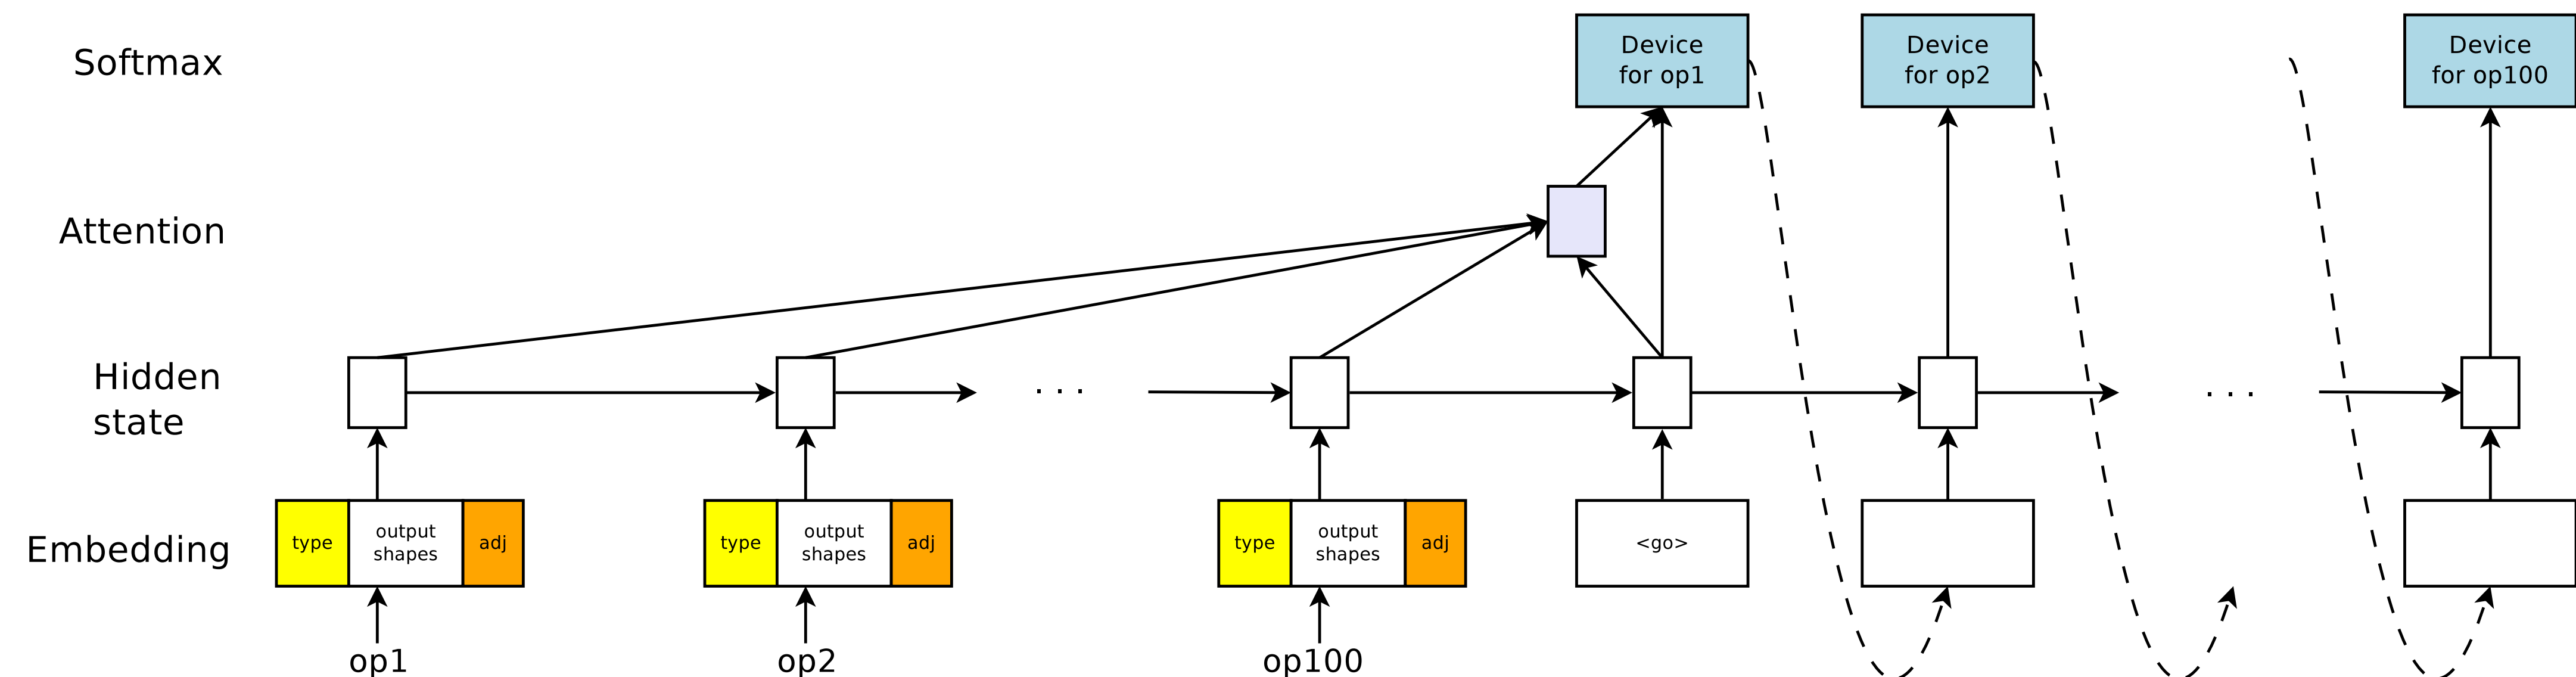
Kullback-Leibler (K-L) divergence



source: Wikipedia

**Back to our device placement problem  
using deep reinforcement learning**

# ICML 2017: a variant of REINFORCE policy gradient



**Neural network model:** a sequence-to-sequence model with an attention layer

**Reward:** square root of the average running time over several iterations


# Our recent work: PPO + cross-entropy

**Reward:**  $\bar{R} - R$

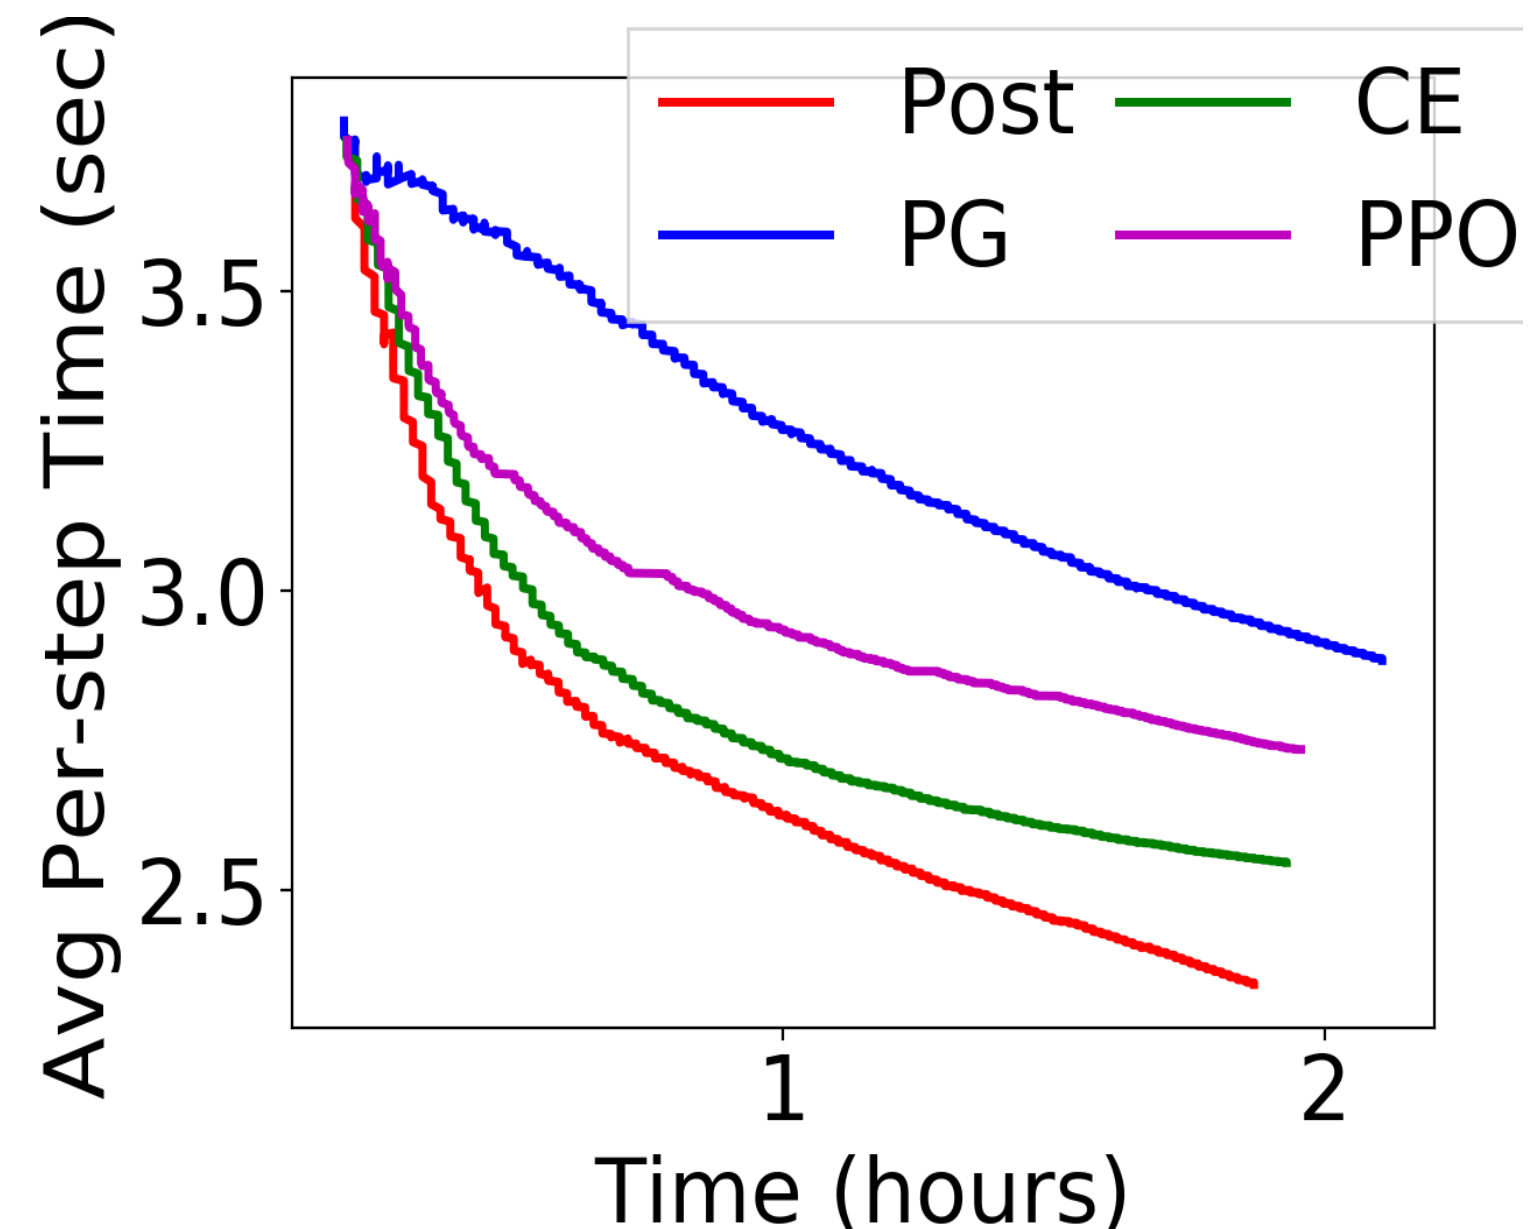
**Every  $K$  steps:** update the policy using PPO

**Every  $N$  steps:** update the policy using cross-entropy


# Training progress over time



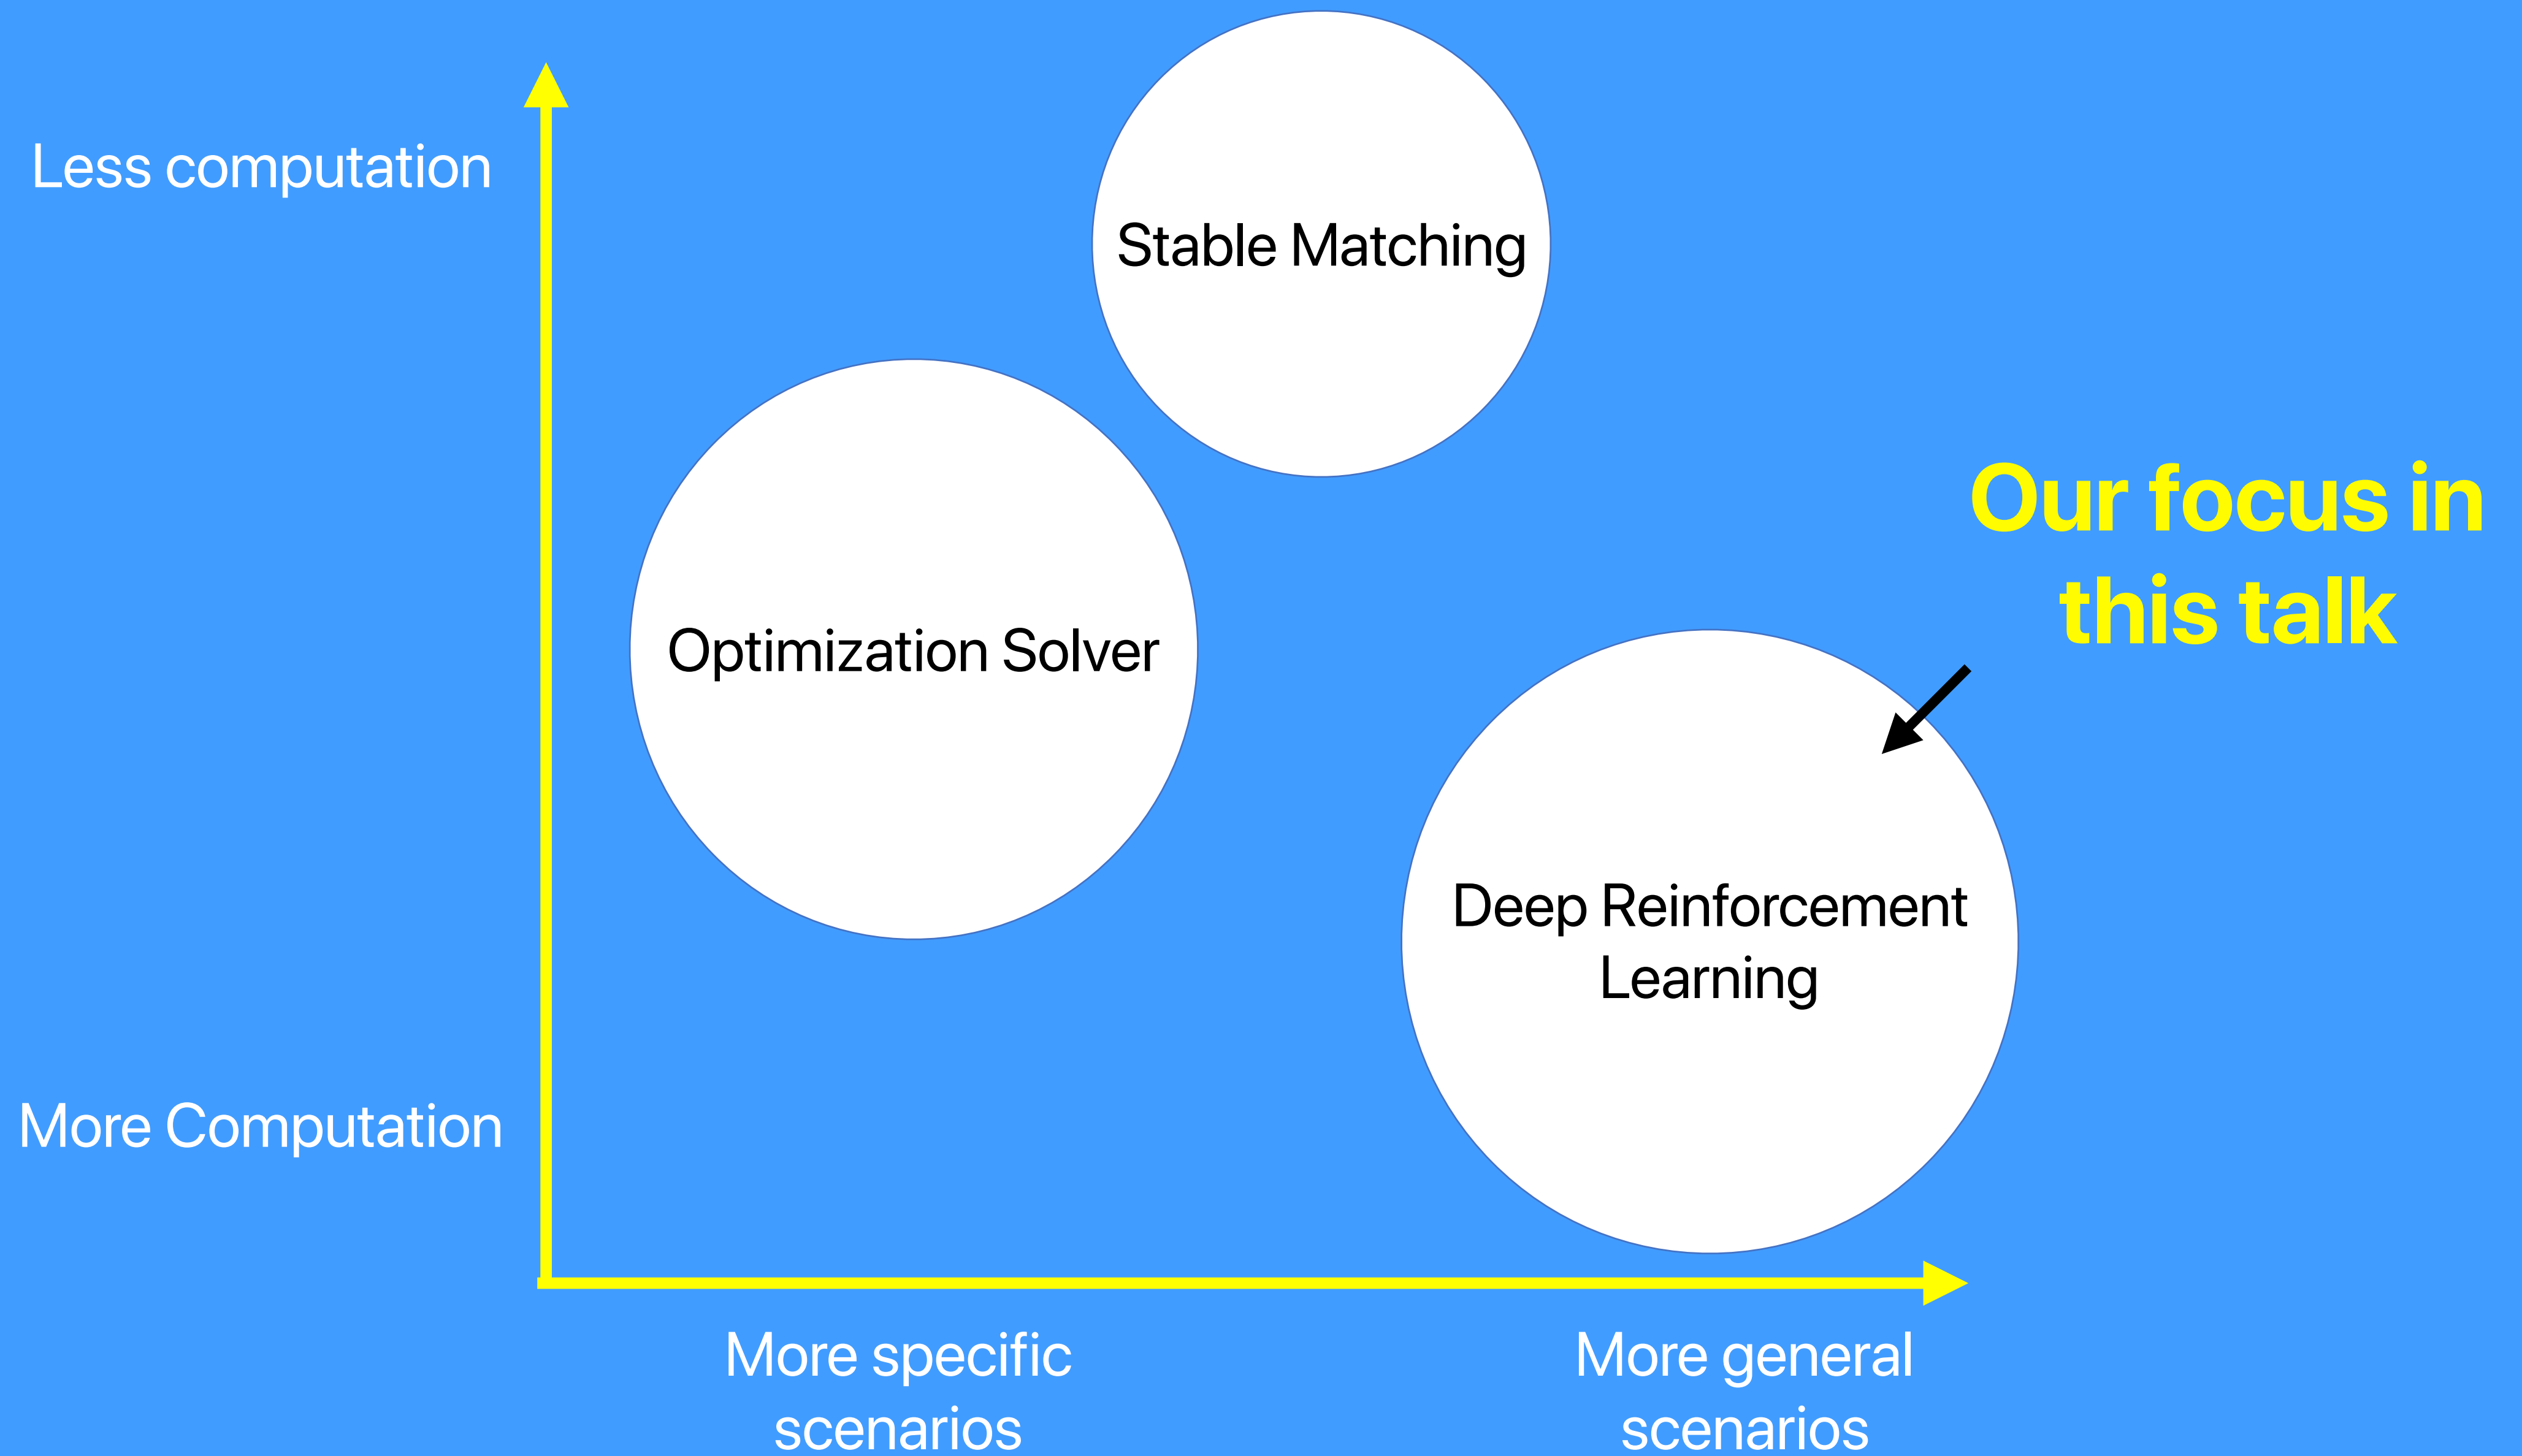
(a) ResNet-50 on 2 GPUs



(b) Inception-V3 on 4 GPUs



(c) ResNet-101 on 8 GPUs



Deep RL may perform quite well in **resource assignment** problems

# Using DRL to solve real-world problems

- ▶ Using DRL for traffic optimization in datacenter networks
- ▶ Congestion control
- ▶ Scheduling distributed machine learning workloads in clusters
- ▶ Federated learning



# Problems with reinforcement learning

- ▶ **Sample inefficiency:** learning a policy usually needs more samples than you think it will
- ▶ **It's difficult to design the reward function**

# Slides and source code:

[iqua.ece.toronto.edu/~bli/iwqos-talk-baochunli.zip](http://iqua.ece.toronto.edu/~bli/iwqos-talk-baochunli.zip)

Baochun Li

Department of Electrical and Computer Engineering

University of Toronto

IWQoS 2019, Phoenix, Arizona

June 24, 2019